

# COMP540 Statistical Machine Learning

Spring 2020  
HW2

Yunhao Zheng (yz157)  
Ziqing Dai (zd15)

# 1. Gradient and Hessian of J() for logistic regression (20 points)

$$\begin{aligned}
 1.1 \cdot g(z) &= \frac{1}{1+e^{-z}} \\
 \frac{\partial g(z)}{\partial z} &= -\frac{1}{(1+e^{-z})^2} \cdot (e^{-z}) \cdot (-1) \\
 &= \frac{e^{-z}}{(1+e^{-z})^2} \\
 &= \left( \frac{1}{1+e^{-z}} \right) \cdot \left( \frac{e^{-z}}{1+e^{-z}} \right) \\
 &= \frac{1}{1+e^{-z}} \left( 1 - \frac{1}{1+e^{-z}} \right) \\
 &= g(z) \cdot (1-g(z))
 \end{aligned}$$

$$\begin{aligned}
 1.2 \cdot J(\theta) &= -\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log(h_{\theta}(x^{(i)})) + (1-y^{(i)}) \log(1-h_{\theta}(x^{(i)}))) \\
 &\quad + \frac{\lambda}{2m} \sum_{j=1}^d \theta_j^2 \\
 h_{\theta} &= g(\theta^T x) = \frac{1}{1+e^{-(\theta^T x)}} \\
 \therefore \frac{\partial J(\theta)}{\partial \theta} &= -\frac{1}{m} \sum_{i=1}^m \left[ y^{(i)} \cdot \frac{e^{-(\theta^T x)}}{1+e^{-(\theta^T x)}} - (1-y^{(i)}) \left( \frac{1}{1+e^{-(\theta^T x)}} \right) \right] x^{(i)} \\
 &\quad + \frac{\lambda}{m} \sum_{j=1}^d \theta_j \\
 &= -\frac{1}{m} \sum_{i=1}^m \left[ y^{(i)} \left( \frac{e^{-(\theta^T x)}}{1+e^{-(\theta^T x)}} \right) - \frac{1}{1+e^{-(\theta^T x)}} \right] x^{(i)} + \frac{\lambda}{m} \sum_{j=1}^d \theta_j \\
 &= -\frac{1}{m} \sum_{i=1}^m (y^{(i)} - h_{\theta}(x^{(i)})) x^{(i)} + \frac{\lambda}{m} [\theta_1, \theta_2, \dots, \theta_d]^T
 \end{aligned}$$

$$1.3 \quad \frac{\partial J(\theta)}{\partial \theta} = \frac{1}{m} X^T \cdot (h_{\theta}(x) - y) + \frac{\lambda}{m} \theta^T$$

$$1.4. \quad H = \frac{1}{m} (X^T S X + \lambda I)$$

$$S = \text{diag} (h_{\theta}(x^{(1)}) (1 - h_{\theta}(x^{(1)})), \dots, h_{\theta}(x^{(m)}) (1 - h_{\theta}(x^{(m)})))$$

$$H = \frac{1}{m} (X^T S X) + \frac{1}{m} \lambda I$$

now, we need to proof

$z^T H z > 0$  for any non-zero  $z$ .

$$z^T H z = \frac{1}{m} z^T X^T S X z + \frac{1}{m} \lambda I z$$

$$= \frac{1}{m} a^T S a + \frac{1}{m} z^T \lambda I z$$

because  $\lambda > 0$ ,  $\frac{1}{m} z^T \lambda I z > 0$

now, we need to proof  $\frac{1}{m} a^T S a > 0$

Because  $z$  is nonzero vector and  $X$  is full rank  
so  $Xz$  is non-zero vector.

$$a^T s a = a_1^2 h_{\theta}(x^{(1)}) (1 - h_{\theta}(x^{(1)})) + \dots + a_m^2 h_{\theta}(x^{(m)}) (1 - h_{\theta}(x^{(m)}))$$

$$a_i^2 > 0 \text{ for } i=1, 2, \dots, m$$

$$\text{since } 0 < h_{\theta}(x^{(i)}) < 1$$

$$\Rightarrow (1 - h_{\theta}(x^{(i)})) > 0$$

$$\Rightarrow z^T H z > 0$$

Therefore,  $H$  is positive definite.

1.5 • Newton's method :

$$\theta_{t+1} = \theta_t - H^{-1} \nabla_{\theta} J(\theta)$$

### Python script

```
# newton's method
```

```
def newton(theta):
```

```
    h = 1/(1 + np.exp(-np.matmul(X, theta)))
```

```
#    J = np.sum(-y * np.log(h) - (1-y) * np.log(1-h)) / m
```

```

# print(J)
grad = np.zeros((dim,))
grad[0] = np.sum(X[:, 0] * (h - y), 0) / m
grad[1:] = (np.sum(X[:, 1:] * (h - y)[:, np.newaxis], 0) + reg * theta[1:]) / m
S = np.zeros((4, 4))
np.fill_diagonal(S, h*(1-h))
Hessian = (np.matmul(np.matmul(X.T, S), X) + reg*np.identity(3))/m
theta = theta - np.matmul(np.linalg.inv(Hessian), grad.T)
print(theta)
return theta

```

```

X = np.array([[1, 0, 3],
              [1, 1, 3],
              [1, 0, 1],
              [1, 1, 1]])
y = np.array([1, 1, 0, 0])
theta = np.array([0, -2, 1])
reg = 0.07
m = 4
dim = 3

```

```

theta = newton(theta)
theta = newton(theta)

```

### Result

```

theta1 = [-3.15199171 -0.40585887  1.81504991]
theta2 = [-4.26505811 -0.29747087  2.33806757]

```

## 2. Overfitting and unregularized logistic regression

Show that for a linearly separable dataset, the maximum likelihood solution for the logistic regression model is obtained by finding a parameter vector whose decision boundary  $\theta^T x = 0$  separates the classes and then, by taking the magnitude of  $\theta$  to infinity. What does this result physically mean? How can we avoid this singular solution?

$$\begin{aligned}\log P(D|\theta) &= \sum_{i=1}^m y^{(i)} \log \frac{1}{1+e^{-\theta^T x^{(i)}}} + (1-y^{(i)}) \log \left(1 - \frac{1}{1+e^{-\theta^T x^{(i)}}}\right) \\&= -\sum_{i=1}^m y^{(i)} \log(1+e^{-\theta^T x^{(i)}}) + (1-y^{(i)}) \log(1+e^{\theta^T x^{(i)}}) \\&= -\sum_{i=1}^m \log(1+e^{-y^{(i)} \theta^T x^{(i)}})\end{aligned}$$

Data is linear separable  $\Rightarrow$  There is  $\theta$  to satisfy  $y \theta^T x > 0$ .

To achieve the maximum of log likelihood, we need to achieve the minimum of  $1+e^{-y \theta^T x}$ , ~~which is~~ which is the maximum of  $y \theta^T x$ .

$$y \theta^T x = \|y\| \cdot \|\theta\| \cdot \|x\| \cos \angle \theta, x.$$

$y, x$  are data. The angle between  $\theta, x$  determines the hyperplane we found.

$\Rightarrow$  We can take the magnitude of  $\theta$  to infinity to achieve maximum log likelihood. It means that after we found the hyperplane to separate data, the ~~loss~~ loss will keep decreasing but we just keep expanding  $\|\theta\|$ .

To ~~avoid~~ avoid this situation: add regularization to our model to keep  $\|\theta\|$  from going to infinity.



### 3. Implementing a k-nearest-neighbor classifier

#### Problem 3.1 Distance matrix computation with two loops (5 points)

See code in k nearest neighbor.py and knn.ipynb

#### Problem 3.2 Compute majority label (5 points)

When  $k=1$ , accuracy= 0.274000

When  $k=5$ , accuracy=0.278000

#### Problem 3.3 Distance matrix computation with one loop (5 points)

See code in compute distances one loop in k nearest neighbor.py

#### Problem 3.4 Distance matrix computation with no loops (5 points)

Two loop version took 18.779556 seconds

One loop version took 25.988581 seconds

No loop version took 0.118588 seconds

#### Problem 3.5 Choosing k by cross validation (5 points)




Figure 1. Cross validation on k


Best result  $K=10$ ,

Got 141 / 500 correct => accuracy: 0.282000

## 4 Implementing logistic regression (45 points)

### Problem 4A1: Implementing logistic regression: the sigmoid function (5 points)

### Problem 4A2: Cost function and gradient of logistic regression (5 points)




### Problem 4A3: Prediction using a logistic regression model (5 points)

Accuracy on the training set = 0.8900

## Problem 4, Part B: Regularized logistic regression (20 points)

### Problem 4B1: Cost function and gradient for regularized logistic regression (10 points)





### Problem 4B2: Prediction using the model (2 points)

Accuracy on the training set = 0.8305

### Problem 4B3: Varying $\lambda$




Figure 2. Decision boundary for  $\lambda = 0$






Figure 3. Decision boundary for  $\lambda = 100.0$

## Problem 4B4: Exploring L1 and L2 penalized logistic regression


### 1. L2 regularization




(a)  $\lambda = 0.3$ : loss = 0.3946  
# Non-zero coefficients = 28



(b)  $\lambda = 1.0$ : loss = 0.4684  
# Non-zero coefficients = 28





(c)  $\lambda = 3.0$ : loss = 0.5478  
# Non-zero coefficients = 28

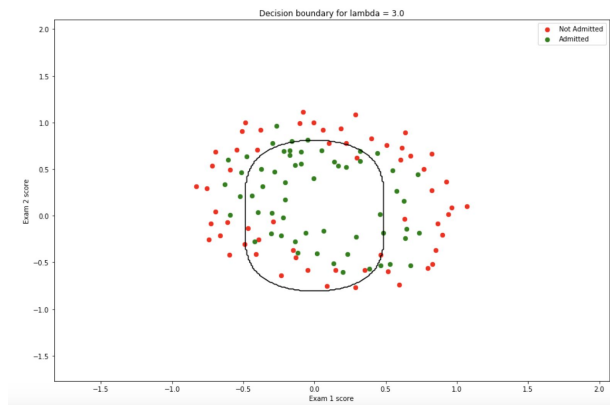


(d) learning path


### 2. L1 regularization



(a)  $\lambda = 0.3$ : loss = 0.3573  
# Non-zero coefficients = 8



(b)  $\lambda = 1.0$ : loss = 0.4381  
# Non-zero coefficients = 7



(c)  $\lambda = 3.0$ : loss = 0.6137  
# Non-zero coefficients = 3

(d) learning path

From the learning path, we see that the coefficients of L1 regularized model shrink faster than that of L2 regularized model as  $\lambda$  (i.e.  $1/C$ ) increases. When  $\lambda$  is large, the L1 regularization provides larger penalty (loss) than L2 regularization, which results in less non-zero coefficients and a simpler model.

#### Problem 4 Part C: Logistic regression for spam classification

Fitting regularized logistic regression models (L2 and L1)

##### L2 penalty

- Standardize features  
Accuracy = 0.9219  
# Non-zero coefficients = 58
- Log transform features  
Accuracy = 0.9434  
# Non-zero coefficients = 58
- Binarize features  
Accuracy = 0.9284  
# Non-zero coefficients = 58

##### L1 penalty

- Standardize features  
Accuracy = 0.9225  
# Non-zero coefficients = 52
- Log transform features

Accuracy = 0.9453

# Non-zero coefficients = 49

c. Binarize features

Accuracy = 0.9284

# Non-zero coefficients = 48

L1 penalty results in more sparse models (model with less non-zero coefficients). I will use model trained by log transform features with L1 penalty because it produces a simpler model with the best accuracy.